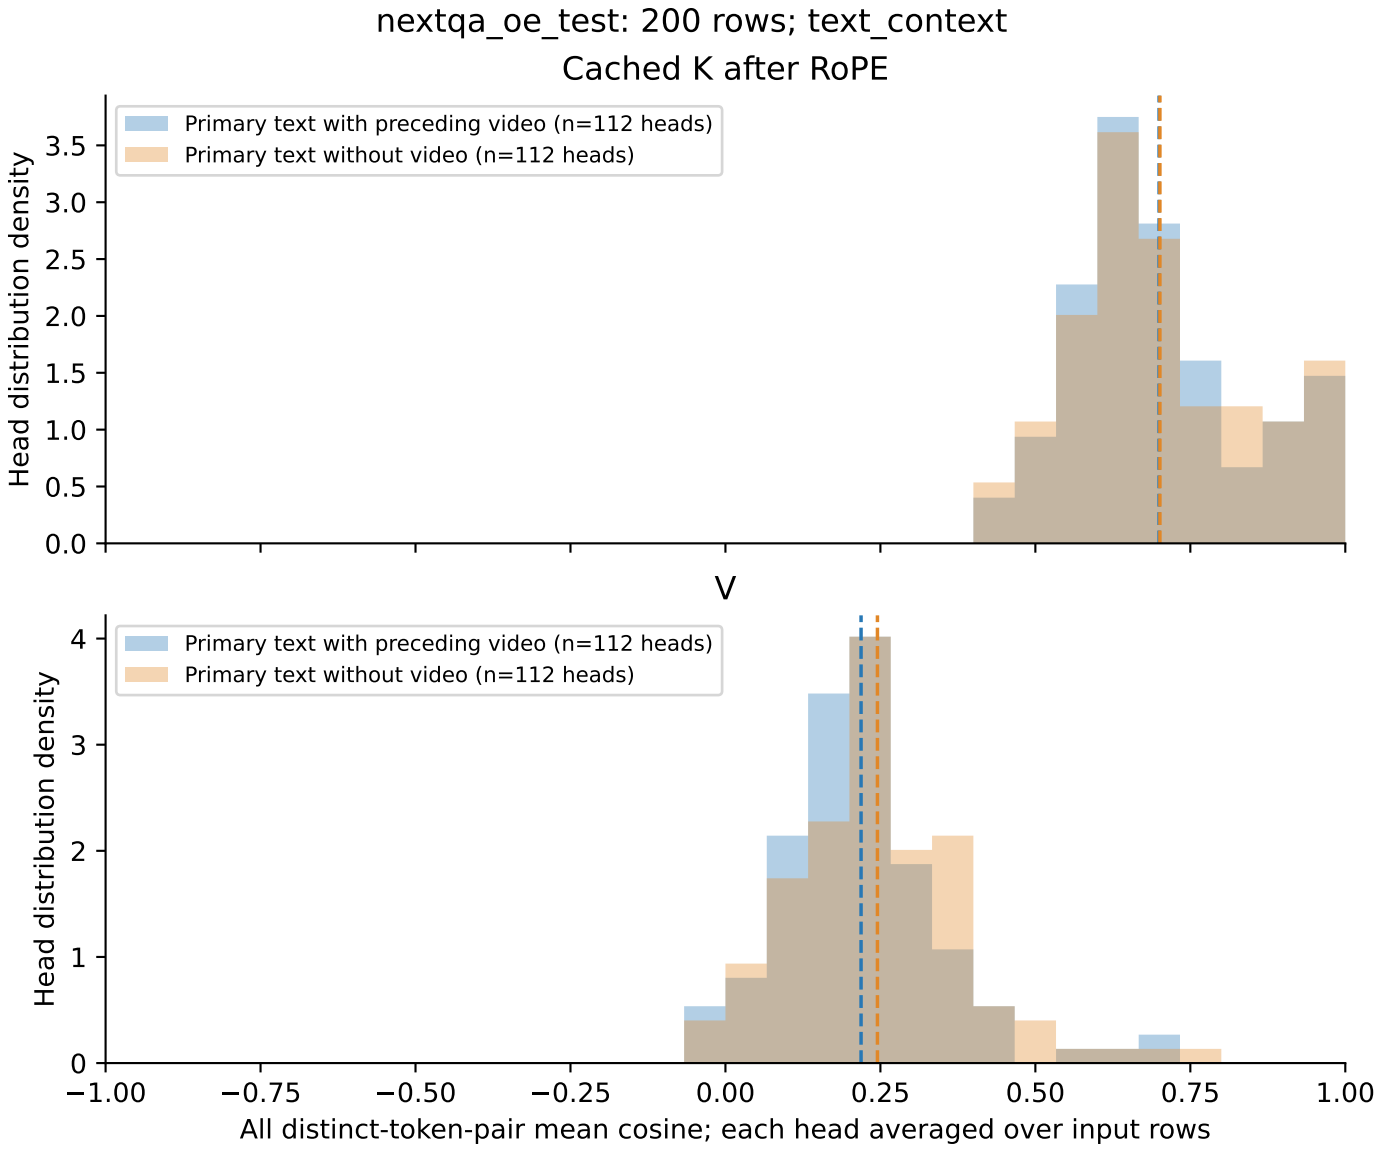
nextqa_oe_test: 200 rows; text_context
Cached K after RoPE
Primary text with preceding video (n=112 heads)
Primary text without video (n=112 heads)
Head distribution density
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
V
Primary text with preceding video (n=112 heads)
Primary text without video (n=112 heads)
Head distribution density
0
1
2
3
4
−1.00
−0.75
−0.50
−0.25
0.00
0.25
0.50
0.75
1.00
All distinct-token-pair mean cosine; each head averaged over input rows